
High-Paying Data Engineer Playbook

5 years experience. Current 13 LPA. Target first switch 28–38 LPA, second jump 40–55 LPA. Yeh beginner course nahi hai — salary-correction + senior positioning ka end-to-end plan hai.

PROFILE

Data Engineer · 5 YOE

STARTING POINT

13 LPA · underpaid vs market

6-MONTH TARGET

28–38 LPA product / GCC

18-MONTH TARGET

40–55 LPA senior / staff-track

5 saal pe 13 LPA ka matlab skill zero nahi hai. Matlab galat employer type + shallow production stories. Big companies (Flipkart, PhonePe, Walmart, JPMC, Google, Amazon, Databricks-class) 5 YOE ko **Spark tuner + pipeline owner + system designer** ki tarah interview karti hain, ETL operator ki tarah nahi.

CONTENTS

Kya-kya cover hai

01 5 YOE pe asli diagnosis

13 LPA kyun stuck hai, market kya pay karti hai, kya skip karna hai

02 Target companies + salary map

Product, GCC, FAANG-track, analytics, services — kahan apply karna hai

03 Senior Data Engineer profile (paper pe kaise dikhna hai)

Title, headline, skill grouping, 5-year story

04 Chronological skill order (5 YOE compressed)

6 months, 10 skills, beginner Hadoop path nahi

05 Month-by-month sprint + resources

SQL expert → Spark internals → dbt/Airflow → cloud/Databricks → Kafka → design

06 3 flagship projects (interview converters)

Batch Spark, warehouse+dbt, CDC streaming lakehouse

07 Resume + LinkedIn rewrite (5 YOE)

Metrics, ownership, ATS keywords, message scripts

08 Interview playbook — big company rounds

SQL, Python, Spark, Airflow, Kafka, system design, HM round

09 90-day calendar + apply engine

Weekly hours, 40 applications, referral system

10 Offer negotiation + checklists

28 vs 32 vs 40 LPA scripts, final packing list

Is PDF ko print karke wall pe phase checklist chipkao. Courses collect mat karo. Har phase ka ek GitHub repo + ek interview story mandatory hai.

5 YOE pe asli diagnosis

5 years experience par 13 LPA India 2026 mein **market-below** hai. Yeh playbook “naya seekho” nahi, **underpaid senior ko correctly pack karke product/GCC mein daalna** hai.

13 LPA

Abhi · typical services/
captive band

16–26

3–5 YOE product mid band

28–45

5–8 YOE product/GCC
senior

40–70

Lead / platform / top-tier
18 months out

Market fact (2026 India)

Experience	IT services / body-shop	Product / unicorn	GCC (Walmart, JPMC, Target...)
3–5 years	8–15 LPA	16–28 LPA	22–38 LPA
5–8 years (senior IC)	15–25 LPA	28–45 LPA	32–55 LPA
Kafka + Spark tuning + Databricks	+2–4 LPA	+5–10 LPA	+5–12 LPA

Bands typical CTC, city Bangalore/Hyderabad/Pune. Loyalty hike services ke andar 8–15% — gap close nahi hota. Switch hi multiplier hai.

SEEDHA BOLUN

5 saal · 13 LPA almost always in 3 buckets: (1) TCS / Infosys / Wipro / Cognizant / Capgemini / Accenture data practice, (2) analytics consultancy (Tiger, LatentView, Fractal-type) with reporting-heavy work, (3) captive team jahan ADF/SSIS/Informatica/Glue jobs run karte ho lekin Spark internals, cost, SLA, on-call ownership nahi. Skills “hai” — **packaging + depth + employer type** nahi.

5 YOE interviewer kya assume karta hai

Big company expect karti hai

- SQL expert — windows, plans, incremental loads
- PySpark production — shuffle, skew, AQE, file sizes
- Ek cloud deep (IAM, storage, cost, networking basics)
- Orchestration: Airflow / Composer / ADF + *idempotency*
- Data modeling: grain, SCD2, late data
- Ek failure story: "job toot gayi, maine yeh fix kiya"
- System design: 100M events/day pipeline on whiteboard
- Stakeholders: analytics/ML ko SLA pe data diya

13 LPA profile often actually hai

- SQL chalta hai, explain plan weak
- PySpark notebooks / Glue jobs, no tuning write-up
- Cloud = "S3 pe file daali" / ADF drag-drop
- Informatica / SSIS / Talend years — keyword trap
- No Git PR culture, no tests, no CI
- Volume unknown ("bahut tables hain")
- Design round pe tools list, trade-offs nahi
- Resume: 20 tools, 0 metrics

Yeh 5 YOE playbook kya NAHI hai

- Python for beginners, Hadoop 2015 path, 12 certifications race
- AWS + Azure + GCP teeno "cover" karna
- Data analyst → DE career switch syllabus
- 2 saal wait karke "in-company promotion se 18 LPA"

EK RULE

Agle 6 months ka job: **3 production-style repos + Spark internals fluency + 1 cloud lakehouse + 8 STAR stories + 40 targeted applications.** Certificate optional. Referral + design round = offer.

Target companies + kahan apply karna hai

5 YOE par shotgun Naukri blast mat karo. **Employer type choose karo** — yahi CTC decide karta hai.

Apply priority (upar se neeche)

Tier	Examples	5 YOE typical	Interview style
A. Product / fintech / consumer	PhonePe, Razorpay, CRED, Flipkart, Swiggy, Zomato, Meesho, Groww, Zerodha, Dream11, ShareChat, Zepto, Blinkit, Paytm	28–42 LPA	Hard SQL + Spark + design + coding
B. GCC / global product India	Walmart Global Tech, Target, JPMC, Goldman, Barclays, Wells Fargo, Amex, Uber, Microsoft, Amazon, Google, Atlassian, Salesforce, Oracle, SAP, Optum, Thomson Reuters	30–50 LPA (Google/Amazon/MSFT upper)	SQL + Python + Spark + design + behavioral
C. Data platforms	Databricks, Snowflake partners, ThoughtSpot, Thoughtworks data, LatentView (only senior IC), Affine (skip if low band)	28–55 LPA	Platform depth, Delta/Iceberg, cost
D. Strong mid product	Freshworks, Chargebee, Postman, BrowserStack, IDfy, Darwinbox, Whatfix, Innovaccer	24–36 LPA	Modern stack, dbt + cloud
E. Avoid as “upgrade”	TCS/Infosys/Wipro/Cognizant/HCL volume hiring, 15% hike wale internal moves	14–20 LPA	Time waste for your goal

5 YOE SWEET SPOT

Pehla switch: **Tier A or B, 28–38 LPA**. “Dream FAANG only” mat baitho 9 months. 2–3 offers collect karo. Google/Amazon parallel try. Services counter-offer almost never match product band.

Roles ke titles — search yeh karo

- **Primary:** Data Engineer, Senior Data Engineer, Data Platform Engineer
- **Also:** Analytics Engineer (dbt-heavy, often 24–34 LPA), Spark Engineer, Databricks Engineer
- **Avoid if you want DE pay:** BI Developer, ETL Developer, MIS, “Data Analyst + SQL”
- **Stretch (month 12+):** Staff Data Engineer, Streaming Engineer, Data Architect (only if design round strong)

City strategy

Bangalore / Hyderabad / Pune / Gurgaon max density. Remote India product possible. US remote 5 YOE possible after 1 strong product stint — pehle India correction, phir dollar rate.

Senior Data Engineer profile — paper pe kaise dikhna hai

Recruiter 10 second mein decide karti hai. 5 YOY ko “learning Spark” nahi likhna. **Owning pipelines at scale** likhna hai — numbers current job se nikaalo, even if company services hai.

Positioning statement (copy, edit, use everywhere)

```
Senior Data Engineer | 5 years building production data platforms
Spark · SQL · Python · [AWS or Azure] · Databricks/Delta · Airflow · dbt
I design, build and operate batch + incremental pipelines with SLAs, data quality
and cost control – from ingestion to modeled warehouse / lakehouse tables
used by analytics and downstream apps.
```

LinkedIn headline (pick one)

- Data Engineer (5 YOY) | PySpark · SQL · Databricks · Airflow | Building reliable lakehouse pipelines
- Senior Data Engineer | Spark performance · dbt · Kafka | Ex-[Company] | Open to Product / GCC

Skill block — recruiter/ATS ke liye (group karke)

Group	Write this (honest only)	Mat likho agar fake hai
Languages	SQL (expert), Python, PySpark	Scala, Java — unless real
Processing	Apache Spark, Spark SQL, Structured Streaming (if true)	Hadoop, MapReduce, Hive as headline
Orchestration	Airflow / Azure Data Factory / Glue Workflows (jo actual hai)	10 orchestrators
Transform / quality	dbt, Great Expectations / dbt tests	Monte Carlo if never used
Storage	Delta Lake / Iceberg, Parquet, S3/ADLS, Snowflake or BigQuery or Redshift or Synapse	Every warehouse logo
Streaming	Kafka, CDC (Debezium/DMS) — after Phase 5	Flink + Pulsar + Kinesis saath
Cloud	One: AWS or Azure or GCP + IAM + cost basics	Multi-cloud expert
DevEx	Git, CI/CD, Docker, Terraform (bonus)	K8s guru unless true

5-year story — 1 paragraph (HM round)

```
I spent 5 years in data engineering, mostly [batch ETL / warehouse / reports]
for [domain: banking / retail / telecom]. I owned [X pipelines / Y tables]
serving [Z users]. The gap I am closing now is production Spark performance,
lakehouse modeling (Delta + dbt), and streaming/CDC – and I want to do that
inside a product or GCC team where data is a core platform, not a ticket queue.
```

Honesty plus ambition. "Mera previous company kharab thi" mat bolna. "Platform ownership chahiye" bolna.

Numbers nikaalo current job se (mandatory homework, week 1)

- ☐ Daily/weekly data volume: GB or billion rows
- ☐ Number of pipelines you touch; which ones you *own*
- ☐ SLA: dashboard by 8am? freshness 1 hour?
- ☐ Worst incident: duration, root cause, your fix
- ☐ Any runtime drop or cost save (even 20% counts)
- ☐ Downstream: finance MIS, ML features, product analytics?
- ☐ Stack exact names + versions if possible

Agar volume nahi pata — Spark UI / job history / warehouse query history / S3 size se nikaalo. "Pata nahi" 5 YOE pe red flag hai.

Services language → product/GCC language (5 YOE translation)

Same work, different words. Recruiter ATS aur HM dono yahi filter karte hain.

Mat likho (services)	Likho (senior DE)
Worked on multiple client requirements in Agile	Owned daily ELT for [domain] used by [N] analysts; SLA 07:00
Developed mappings in Informatica / ADF pipelines	Built incremental Spark/dbt models (MERGE) on [cloud] with tests
Involved in data extraction from source systems	Ingested CDC/batch from [Oracle/SQL Server/API] into bronze Delta/S3
Created reports for stakeholders	Published gold marts (grain: 1 row per order) consumed by Looker/Power BI
Handled production issues	On-call: cut MTTR from [x]h to [y]m with alerting + runbooks
Optimized queries as needed	Tuned Spark job via UI: shuffle [A]→[B] GB, runtime [T1]→[T2]
Knowledge of Hadoop ecosystem	Delete. Write Spark/Delta/cloud you actually ran.

Chronological skill order — 5 YOE compressed

Beginner 12-month path mat follow karo. 5 saal baad order yeh hai: **depth pehle, naya tool baad mein.**

#	Skill	When	5 YOE bar (itna hi kaafi nahi)	Interview signal
1	SQL expert + query plans	Month 1	Windows, gaps-and-islands, incremental MERGE, partition pruning	Every company, round 1
2	Python for pipelines	Month 1	typing, pytest, retries, dataclasses, not pandas-only	Product coding round
3	Data modeling	Month 1–2	Grain, star, SCD2, late-arriving facts, conformed dimensions	Design + dbt round
4	Spark internals + tuning	Month 2–3	Shuffle, skew, AQE, broadcast, file sizing, spill, UI	Biggest salary lever
5	Airflow (or equivalent) production	Month 3	Idempotency, backfill, SLAs, pools, sensors, data-aware scheduling	GCC + product
6	dbt + tests + contracts	Month 3–4	Incremental, snapshots, exposures, CI	Modern product stack
7	One cloud + lakehouse	Month 4–5	Medallion, IAM, cost, Databricks Jobs or Snowflake tasks	JD keyword #1 after Spark
8	Kafka + CDC	Month 5	Offsets, partitions, exactly-once, late data, when NOT to stream	28%+ pay premium
9	Data system design	Month 5–6	SLA, replay, DQ, cost, CDC vs batch, Iceberg/Delta	Senior bar
10	AI-adjacent data (optional premium)	Month 6+	Feature tables, embedding/RAG ingestion, unstructured → lakehouse	Differentiator, not core

SKIP LIST (5 YOE TIMEPASS)

MapReduce, Pig, Oozie, Tableau-first path, Power BI certification as DE headline, learning 3 clouds, Kubernetes-from-zero unless platform team, Kaggle as DE proof, 6 Udemy certificates without GitHub.

Cloud pick (ek, deep)

IF JDS AROUND YOU ARE AZURE / BANKS / GCC

Azure + Databricks

ADLS Gen2, ADF or Airflow, Databricks (Delta, Unity Catalog, Jobs), Synapse/Snowflake optional. Cert: DP-203 after project, not before.

IF STARTUPS / PRODUCT / AMAZON-LIKE

AWS + Databricks or Glue/EMR + Athena

S3, IAM, Glue or EMR/Spark, MWAA/Airflow, Redshift or Snowflake. Cert: AWS Data Engineer Associate after one deployed pipeline.

GCP only if current job already GCP/BigQuery heavy. Nayi cloud mat add karo “resume pe accha lage”. 5 YOE pe depth jeet-ti hai.

6-month sprint — month by month + resources

Job ke saath **10–12 hours / week**. Weekend pe 5–6 hours project. Weeknights 60–75 min SQL/Spark/design. Certificate last, GitHub first.

WEEKS 1–4 · NO NEW SHINY TOOLS

Month 1 — Foundation reset (senior bar)

Goal: SQL + Python + modeling itne tight ho ki product round 1 clear. Resume v1 live. Numbers homework complete.

SQL — expert checklist

- Window functions: RANK/DENSE_RANK, LAG/LEAD, running totals, sessionization
- Gaps-and-islands, top-N per group, dedupe, SCD2 query patterns
- JOIN types including anti-join / semi-join; EXISTS vs IN
- Incremental load: MERGE / INSERT OVERWRITE partitions
- Tuning: EXPLAIN, partition pruning, clustering/clustering keys, file skip

Python — pipeline, not notebook

- Functions, pathlib, logging, retries (tenacity), argparse
- pytest for transforms; type hints; virtualenv
- Small API ingest script with backoff + schema check

Resources

- **Mode SQL Tutorial** — rewrite mentally as “I already know this, speed-run”
- **Select Star SQL** — windows
- **LeetCode Database** — 80–100 Q, hard included (5 YOE bar)
- **Use The Index, Luke** — plans / indexes
- Book: *The Data Warehouse Toolkit* (Kimball) — star / SCD chapters only
- **pytest docs** + **Real Python — logging, pathlib**
- Practice warehouse: **SQLPad** / **HackerRank SQL advanced**

Milestone: 40 LeetCode DB done · 1-page modeling of your current job’s core fact · Resume v1 on LinkedIn.

Month 2–3 — Spark performance (salary engine)

Goal: “I can read Spark UI and fix a slow job” — yeh sentence interview mein proof ke saath. Syntax wale DE 18 LPA. Tuner 30+.

Deep list (is sequence mein)

1. Driver vs executor, cores, partitions, tasks, stages
2. Narrow vs wide; shuffle write/read; exchange in plan
3. Parquet layout, 128–256 MB files, too-small files problem
4. Partitioning vs bucketing; when not to partition on high-cardinality
5. Joins: broadcast vs sort-merge; skew (salting, AQE skew join)
6. AQE, adaptive coalescing, dynamic partition pruning
7. Caching vs checkpoint; spill to disk; OOM patterns
8. Delta Lake: MERGE, Z-order/OPTIMIZE, vacuum, time travel

Resources

- [Spark SQL programming guide](#)
- Book: *Learning Spark* (O'Reilly, 2nd ed.) — chapters on SQL, joins, tuning
- [Databricks Spark optimization guide](#)
- [Delta Lake docs](#)
- YouTube: Spark UI walkthroughs (Databricks official + SparkByExamples for API only)
- Dataset: NYC TLC trip data / GitHub archive — 10GB+ compulsory

Milestone — Project A started: 10GB+ batch job. Before/after: runtime + shuffle bytes + cost note in README.

Month 3–4 — Airflow + dbt + quality

Goal: Ticket-runner nahi, *data product owner*. Tests, docs, backfill, idempotent DAG.

- **Airflow:** DAG design, task vs DAG failure, retries, pools, sensors, datasets (data-aware), backfill, catchup off, idempotent tasks, Slack/email on fail
- **dbt:** staging/intermediate/marts, ref/source, incremental merge, snapshots (SCD2), tests (unique, not_null, accepted_values, relationships), docs generate, slim CI
- **Quality:** freshness + volume + null spikes. dbt tests pehle; Great Expectations/Soda extra

Resources

- **Astronomer Academy** (Airflow fundamentals + DAG writing) — free
- **Apache Airflow docs** — especially best practices + data-aware scheduling
- **dbt Fundamentals + dbt Best Practices**
- **DataTalks.Club DE Zoomcamp** — modules skip karke Airflow/dbt/batch use karo, beginner videos skip
- **Great Expectations** docs — 5 tests on your mart

Milestone — Project B: ingest → bronze → dbt marts (SCD2) → Airflow daily → tests in CI.

Month 4–5 — One cloud lakehouse (Databricks or Snowflake)

Goal: Notebook tourist nahi. Jobs, IAM, medallion, Unity Catalog/RBAC, cost dashboard. India 2026 JDs: **Databricks** hottest, Snowflake close second.

- Medallion: bronze raw, silver cleaned, gold dimensional / business
- Databricks: clusters vs SQL warehouse, Jobs, Delta Live / Lakeflow if available, Asset Bundles + GitHub Actions
- Snowflake path: warehouses, micro-partitions, clustering, tasks+streams, cost per query
- Security: least-privilege IAM, secrets not in notebooks
- Cost: job ₹500 vs ₹50,000 — same data, different plan

Resources

- **Databricks Academy — Lakehouse fundamentals + DE associate path**
- **Databricks Asset Bundles** + GitHub Actions CI
- **Snowflake Hands-on Essentials** (free badge)
- AWS: **Skill Builder data path** · Azure: **DP-203 learn path**
- Well-Architected analytics lens (AWS) / Azure Cloud Adoption data landing zone docs

Cert timing: Project cloud pe chalane ke *baad* Databricks DE Associate ya DP-203 / AWS DEA-C01. 5 YOE pe cert without project = services signal.

Month 5 — Kafka + CDC (premium band)

Goal: Streaming hype nahi — *offsets, partitions, consumer groups, replay, late events, exactly-once sinks*. Senior signal: kab batch kaafi hai.

- Kafka 101 then production: replication, ISR, key-based partitioning
- CDC: Debezium or AWS DMS — DB binlog → Kafka → lakehouse MERGE
- Spark Structured Streaming: watermark, window, checkpoint, trigger
- Failure: consumer lag, poison pill, schema evolution (Schema Registry)

Resources

- **Confluent Kafka 101** (free) then Kafka Streams only if time
- **Kafka official docs** — design section
- **Debezium documentation**
- **Apache Iceberg** + Delta MERGE as sink patterns
- Book (slow read months 1–6): *Designing Data-Intensive Applications* — Ch. 1–3, 8–11

Milestone — Project C: CDC or clickstream → Kafka → Delta/Iceberg → documented late-data + replay.

Month 6 — System design + interview engine

Goal: Whiteboard pe 45 min: “Design Uber trip pipeline / Stripe ledger warehouse / 100M events/day clickstream.” Apply 40 roles. 8 STAR stories. Mocks.

Resources

- DDIA (continue) — replication, partitioning, stream processing chapters
- YouTube: “data engineering system design” (Alex Xu-style but data: batch vs stream, storage, DQ)
- Educative / ByteByteGo data-oriented case studies if budget
- Mocks: Pramp / interview buddy / Exponent DE videos — pause and answer yourself

Milestone: 2–3 onsite loops. Target offer 28–38 LPA. Walk-away below 24 unless ESOPs insane + brand.

3 flagship projects — 5 YOE interview converters

Tutorials clone mat karo. Har repo mein: architecture diagram, SLA, failure modes, backfill, cost, Spark UI screenshot, “what I would do at 10× scale”.

Project A — Batch lakehouse (Spark tuner story)

Story line: “I ingested NYC taxi / e-commerce events (~10–50GB), built bronze-silver-gold, cut job time from X → Y by fixing shuffle, file sizing and skew.”

Stack: PySpark + Delta on Databricks community / local Spark + S3-compatible (MinIO) · GitHub Actions pytest on transforms

Must show: partition columns, OPTIMIZE/compaction, before-after Spark UI, cost estimate table

README sections: Problem · Architecture · Data model · Tuning log · How to run · Next (streaming)

Project B — Warehouse as a product (Airflow + dbt)

Story line: “Daily ELT from public API + Postgres dump into modeled marts with SCD2 customer dimension, tested in CI, SLA 07:00.”

Stack: Airflow (Docker) + dbt Core + Postgres or Snowflake trial + Great Expectations

Must show: lineage screenshot (dbt docs), failed test → pipeline stop, backfill runbook, idempotent DAG

Senior twist: data contract YAML (owner, SLA, grain) for 2 marts

Project C — CDC / streaming (premium)

Story line: “Orders table CDC via Debezium → Kafka → Spark Structured Streaming MERGE into Delta. Late updates 2 hours. Exactly-once via checkpoint + idempotent merge keys.”

Stack: Docker Compose: Postgres + Debezium + Kafka + Spark + MinIO/Delta

Must show: why MERGE keys, watermark choice, what happens if consumer down 3 hours, schema evolution plan

Skip if: Month 6 interviews already coming — Project A+B rock solid better than half C.

PUBLIC GITHUB HYGIENE (5 YOE)

Repo name: `nyc-taxi-delta-lakehouse` not `spark-practice`. Pin architecture PNG. LICENSE. No secrets. English README. LinkedIn pe 3 posts: tuning, SCD2, CDC late-data — each 8–12 lines + diagram.

Resume + LinkedIn rewrite — 5 YOE

1 page (max 2 if 2 companies). PDF. Metrics. Tools grouped. **ETL developer language hatao.**

Header

```
NAME | Bangalore | +91-XXXXXXXXXX | email | LinkedIn | GitHub
Data Engineer · 5 years · PySpark, SQL, Python, Databricks/Spark, Airflow, dbt, AWS/Azure
```

Summary (4 lines, no fluff)

```
Data engineer with 5 years building and operating production ETL/ELT
pipelines on [Spark/cloud]. Owned freshness SLAs for [N] downstream
dashboards/models; reduced [job] runtime by [x%] and improved
reliability via tests, incremental loads and monitoring.
Seeking Senior Data Engineer roles in product / GCC teams.
```

Bullet formula (har line)

Verb + system + scale + method + result

```
Weak: Worked on Spark jobs for sales data in Azure.
Strong: Rebuilt 12 Spark ETL jobs (~180M rows/day) on Databricks Delta;
        cut cluster hours 41% via repartition, broadcast joins and
        file compaction; dashboard SLA 06:30 → 05:10.
```

Experience section — 5 YOE pattern

- Latest role: 5–7 bullets (ownership first)
- Previous: 3–4 bullets
- Agar 5 saal same company: split by project/promotion internally
- Client names if allowed (banks love GCC keywords: payments, ledger, risk)
- Do not list 15 tools in one bullet

Steal-able bullet bank (numbers apne se badlo)

- Designed star-schema model (orders fact + 6 dims, SCD2 customer) used by finance MIS; reduced ad-hoc Excel extracts ~70%.
- Implemented incremental MERGE on [table], dropping full refresh from 4.5h to 38m; enabled same-day restatements.
- Added Airflow SLAs + retries + alerting; dropped silent overnight failures from ~6/month to 1.
- Standardized Parquet file sizes (~128MB) and partition spec (event_date, country); scan cost down [x] %.
- Built dbt test suite (82 tests) in CI; blocked 3 bad loads from reaching gold in first month.
- Documented runbooks for backfill and late-arriving facts; onboarding time for new joiner ~2 weeks → 4 days.

What to delete

- “Team player, hard working, MS Office, Windows”

- Academic projects if 5 YOE unless famous OSS
- Every Hadoop ecosystem logo
- Objective: "To work in a challenging environment..."
- Photo, DOB, father name, religion, passport (India product/GCC)

LinkedIn outreach (short)

Hi [Name] – I'm a data engineer (5 YOE, Spark/SQL/[cloud]) targeting Senior DE roles on your team. I recently published a lakehouse project with Spark tuning notes + dbt tests: [link]. If you're open to a referral for [Job ID], I can send a 5-line fit summary.

20 hiring managers / senior DEs per week. 3 lines. Link project, not resume dump.

Copy-paste 5 YOE resume skeleton (1 page — numbers replace karo)

RAHUL SHARMA · Hyderabad · rahul@email.com · linkedin.com/in/rahul · github.com/rahul
Data Engineer · 5 years · SQL · Python · PySpark · Databricks/Delta · Airflow · dbt · Azure/AWS

SUMMARY

Data engineer with 5 years building production batch ELT on Spark and [Azure/AWS]. Owned freshness SLAs for finance + product analytics (≈[X] GB/day). Reduced a critical Spark job [4.5h → 42m] and added tests so bad data cannot reach gold.

EXPERIENCE

Company A – Data Engineer 2022 – present

- Owned 14 production Spark jobs (≈120–200M rows/day) on Databricks Delta; cut cluster hours 37% via broadcast joins, 128MB file compaction, partition repair.
- Modeled orders star schema (fact + 5 dims, SCD2 customer) used by 18 analysts; replaced ad-hoc SQL extracts with dbt marts + 70 tests in CI.
- Built Airflow DAGs with retries, SLAs and Slack alerts; silent overnight failures ~6/month → 1. Wrote backfill runbook (idempotent partition overwrite).
- Implemented incremental MERGE for [table]; same-day restatement for finance close.

Company B – Data Engineer / ETL 2021 – 2022

- Migrated 9 Informatica workflows to Spark/ADF; validated row counts ±0.01%.
- Landed source extracts to ADLS/S3 Parquet; documented grain and late-data rules.

PROJECTS (GitHub)

- nyc-taxi-delta-lakehouse – 10GB+ medallion; Spark UI before/after + cost note
- orders-dbt-airflow – SCD2 + tested marts + daily DAG

SKILLS

SQL (windows, incremental, plans) · PySpark · Python/pytest · Delta Lake
Airflow · dbt · [Azure: ADF, ADLS, Databricks] · Git · CI

Interview playbook — badi companies kya poochti hain

5 YOE loop typically: recruiter → SQL/Python → Spark/stack → system design → hiring manager.
Services-style “what is RDD” pe mat atko — **trade-offs + production stories**.

Round map

Round	Duration	Bar at 5 YOE
Recruiter	20–30m	Notice, CTC, location, why leaving, stack match. Ask team’s data scale.
SQL live	45–60m	2–3 hard queries + optimize. Talk plan. Don’t stay silent.
Python / coding	45m	Medium: dicts, heaps, parsing logs, not ICPC. Clean functions + tests talk.
Spark / DE deep dive	60m	UI, skew, joins, Delta MERGE, your project grilling.
System design	45–60m	Clarify SLA, grain, volume, fail modes, cost, backfill, DQ.
HM / behavioral	45m	Conflict, ownership, “tell me a miss”, why product/GCC, 5-year arc.

SQL — 5 YOE question types (drill these)

- Latest order per customer; running GMV; WoW growth
- Sessionize events (30-min gap)
- SCD2 point-in-time: “customer segment on order date”
- Find overlapping intervals / consecutive days streaks
- “This query scans 12TB, how do you fix?” (partition, clustering, predicate pushdown, pre-agg)

Practice aloud. Interviewer process maangta hai, sirf final SELECT nahi.

Spark — expected answers (short, exact)

Shuffle

Wide transformation se data exchange across executors. Costly due to disk + network.
Reduce: pre-filter, broadcast small side, partition for join keys, AQE.

Skew

One key explodes a partition. Symptoms: one task 10× longer. Fixes: salting, skew join (AQE), isolate hot key, two-phase agg.

Broadcast join

Small table fits executor memory (and `spark.sql.autoBroadcastJoinThreshold`). Huge table nahi broadcast. After filter, stats matter.

Small files

Thousands of tiny Parquet files = slow listing + tasks. Compact to ~128MB. coalesce vs repartition trade-off (shuffle).

System design template (har question pe yeh order)

1. **Clarify:** sources, volume (peak), latency SLA, retention, consumers (BI vs ML vs product), exactly-once need?
2. **Grain + model:** facts/dims, keys, SCD, late data policy
3. **Batch vs stream vs CDC:** say when batch wins (cheaper, simpler)
4. **Architecture:** ingest → bronze → silver → gold; queue; warehouse; serving
5. **Reliability:** retries, DLQ, idempotent sinks, replay from Kafka/offset or partition reload
6. **DQ + observability:** tests, freshness alerts, lineage, on-call
7. **Cost + security:** compute isolation, partition pruning, PII, IAM
8. **Scale 10×:** what breaks first

Design prompts — weekly practice (write 1-page answers)

- YouTube watch events: 100M/day, 5-min freshness for trending, 24h for finance
- Ecommerce orders + refunds CDC, finance close T+1, inventory near-real-time
- Bank core ledger → warehouse, 7-year retention, PII masking, SCD2 customer
- IoT 10k devices, 1Hz metrics, downsample for ops dashboard
- Migrate Informatica + on-prem Oracle to Databricks in 6 months (phased)

STAR stories — 8 likh ke ratlo (5 YOE gold)

1. Slow Spark job you diagnosed from UI
2. Pipeline silent fail / bad data reached dashboard
3. Ambiguous requirement, you defined grain/SLA
4. Conflict with analyst / PM on “wrong numbers”
5. Backfill of months of data without locking prod
6. Cost explosion you stopped
7. Mentored junior / documented runbook
8. Deadline miss — what you changed next time

Format: Situation (2 lines) · Task · Action (3–5 technical steps) · Result (number).

Sample STAR (ratlo, numbers apne se)

```
Situation: Finance dashboard missed 08:00 SLA 4 times in a month. Job
“fnl_orders_enrich” ran 4.5h, often after 08:20. Cluster cost also spiked.
Task: I owned the job. Need < 60 min, same numbers, no extra budget.
Action: Spark UI showed 1 skewed stage (join on country=IN, 70% rows).
Enabled AQE skew join, broadcast the 80MB dim, wrote 128MB Parquet
instead of 12k tiny files, switched full refresh to MERGE on order_date
(3-day lookback). Added duration + row-count alerts.
Result: 4.5h → 42m. SLA green for 8 weeks. Cluster hours -41%.
This is the story I tell in every Spark round.
```

CTC / NOTICE TRAPS

Recruiter: current 13, expectation **28–35** for product/GCC (range, not exact). “Flexible for the right platform role.” Notice: 30–90 as per bond; joining 60 days typical. 13 ko hide mat karo, *market correction* frame karo: “Band services wala tha; last 6 months Spark/lakehouse depth + ownership stories.”

Python coding — typical DE set

- Parse nested JSON logs → flatten to records
- Top K frequent IPs (heap / Counter)
- Iterator / generator for large file (memory)
- Retry with exponential backoff sketch
- Two-sum style easy-medium if FAANG; else data wrangling

Resource: [LeetCode](#) 40–60 medium (arrays, hash, strings, heap). 5 YOE DE ko CP-god nahi banna — clean code + complexity talk.

Airflow / dbt / Kafka one-liners they love

- **Idempotent DAG:** rerun same day = same tables, no duplicates (partition overwrite or MERGE keys).
- **Catchup vs backfill:** catchup auto-runs missed intervals (dangerous); backfill is controlled.
- **dbt incremental:** unique_key + lookback window; know when to full-refresh.
- **Kafka exactly-once:** producer idempotence + transaction or sink idempotency (MERGE); “at-least-once + dedupe” often practical.
- **Consumer group:** scale = partitions $\geq$ consumers; lag = production KPI.

90-day calendar + apply engine

5 YOE ko 9-month course nahi chahiye. **90 din mein interview-ready + 40 targeted applications.** Projects parallel, not sequential-perfect.

Week	Build	Drill	Career
1	Numbers homework from current job	SQL 7 Q / weekdays	Resume v1, LinkedIn headline, GitHub polish
2	Kimball 1-pager of your domain	SQL hard + 1 Python script + pytest	10 recruiter searches saved (Tier A/B)
3–4	Project A skeleton + 10GB dataset	Spark architecture notes	5 informational chats
5–6	Project A tuning log (before/after)	Spark UI + join strategy drills	Resume v2 with Project A bullets
7–8	Project B Airflow + dbt staging	LeetCode DB 5/week + Python 3/week	10 referrals messages
9–10	dbt marts + SCD2 + tests CI	First 2 system-design writeups	Start applying: 8 jobs/week
11–12	Cloud deploy Project A or B; cert only if extra time	Mock SQL + mock design	8 jobs/week + follow-ups
13+	Project C if interviews not yet; else interview only	STAR stories aloud	Loops, negotiate, don't freeze applying

Weekly hours (job ke saath)

Block	Time	What
Weeknights × 4	70 min	SQL / Spark internals / design prompt
Saturday	4 hrs	Project build
Sunday	3 hrs	Project + LinkedIn + applications
Friday night	30 min	Log: what shipped this week (keep a brag doc)

Application engine (volume with quality)

- Target: **40 applications in weeks 9–14** — only Tier A/B/C/D from Chapter 02
- Naukri/Instahyre/Cutshort/LinkedIn Easy Apply + *always* 1 employee message
- JD se 6 keywords resume pe naturally (Spark, Databricks, Airflow, Python, SQL, AWS/Azure)
- Track spreadsheet: company, role, date, referral, status, next follow-up
- Ghosting normal hai. 7 din baad polite bump.

DON'T WAIT TO BE "100% READY"

Week 9 se apply. Interviews khud bata denge weak area. 5 YOE pe learning in public + live loops > 6 aur months hiding.

Offers, negotiation, checklists

Anchor

Situation	Ask / walk-away
First product/GCC offer, 1 offer only	Ask 32–36 if they opened 28. Walk-away < 24 unless brand + learning insane (Amazon SDE-T type may still be worth it).
2+ offers	Play written competing offer. 10–18% extra common.
Current 13, they lowball 18 “50% hike”	Politely reject. 50% hike services math hai, 5 YOE market nahi.
ESOPs heavy startup 24 + 0.0x%	Discount ESOP. Cash 28+ safer for correction.

Script

I’m excited about the team and the data platform problems. Based on 5 years of production DE and the senior scope we discussed, I was expecting CTC in the 32–36 LPA range. If you can revise the cash component, I’m ready to sign this week.

Joining

Buyout/notice: get offer letter before resign. Don’t resign on verbal. Background verification: keep payslips, increment letters. Non-compete usually weak in India for DE but don’t steal data/code.

18-month second jump (40–55)

- Year 1 in product/GCC: own a domain end-to-end, on-call, cost, 1 streaming or CDC production
- Staff-track signals: design docs, multi-team APIs, SLAs others depend on
- Then: Databricks/Snowflake/FAANG-class or Staff DE internal promo

Final packing list

Ship

- ☐ Resume 1-pager, metrics, PDF
- ☐ LinkedIn headline + 3 project posts
- ☐ GitHub: Project A + B (C optional)
- ☐ 80 SQL questions + 8 STAR stories
- ☐ 5 design one-pagers
- ☐ 40 targeted applications
- ☐ Salary range practiced aloud

Don’t

- 3 clouds, 12 certs, 0 repos
- Services 15% hike as “win”
- Title Junior / Trainee on resume
- Fake Kafka/Databricks
- Silence in SQL round
- Apply only FAANG for 8 months
- Resign before written offer

Resource appendix (all links)

- mode.com/sql-tutorial
- selectstarsql.com
- leetcode.com/problemset/database
- use-the-index-luke.com
- [Spark SQL guide](#)
- docs.delta.io
- [Databricks Academy](#)
- academy.astronomer.io
- [dbt Fundamentals](#)
- [dbt best practices](#)
- [DE Zoomcamp](#)
- [Confluent Kafka 101](#)
- kafka.apache.org/documentation
- debezium.io/documentation
- iceberg.apache.org
- learn.snowflake.com
- skillbuilder.aws
- [Azure DP-203 path](#)
- greatexpectations.io
- Books: Kimball Toolkit · Learning Spark 2e · DDIA

CLOSING

5 years experience tumhara liability nahi, **undervalued asset** hai. Market 5 YOE DE ko 28–45 LPA isliye deti hai kyunki production judgment khareedti hai — tools nahi. Agle 90 din: depth, 2 repos, 8 stories, 40 applications. Badi companies ki “bharmaar” cold PDF se nahi, **targeted loops + proof of Spark/SLA/ownership** se aati hai.

Playbook version 1.0 · Tailored for 5 YOE Data Engineer · Current 13 LPA · India product/GCC track 2026. Numbers are typical bands, not a guarantee. Update resume with only truthful metrics.