

Manish Pandey

Data Engineer · Python · SQL · PySpark · dbt · Azure Databricks · CDC

Gurugram, India · +91-9140357957 · mp4966@gmail.com · 5 years

SUMMARY

Data Engineer with 5 years of production experience building ELT/ETL platforms, warehouses, and CDC pipelines. Currently own the **Staff Service Commission Candidate 360** platform: **dbt** incremental models on **Azure Databricks** (Spark SQL / PySpark, Delta), tested in CI, served on **Apache Superset**. Also shipped production **CDC** (GoldenGate), Spark/SQL performance work, multi-region migrations, and GitLab CI/CD for data releases. Own modeling, incremental loads, data quality, and on-call fixes end to end.

SKILLS

Languages	Python, SQL (windows, MERGE, incremental loads, query plans), PySpark / Spark SQL, PL/SQL
Lakehouse	Azure Databricks, Delta Lake, dbt (staging / intermediate / marts, incremental, tests, docs), Apache Superset
Pipelines	ETL / ELT, CDC, dimensional modeling, Candidate 360, data quality, reconciliation, lineage
CDC & ETL	Oracle GoldenGate (Extract/Replicat, lag, trail, recovery), ODI, Data Guard, Oracle, MariaDB
Engineering	GitLab CI/CD, Git, Azure, production monitoring, incident RCA, runbooks

EXPERIENCE

Data Engineer (Consultant)

Apr 2022 – Present · Gurugram

Cubastion Consulting

Production data platforms for exam analytics, telecom, and automotive — ingest, transform, warehouse, CDC, and serve.

Staff Service Commission (SSC) — Candidate 360 Data Platform **Current** · **dbt** · **Databricks** · **PySpark** · **Superset**

- Built production ELT: sources → **dbt** staging / intermediate / marts, materialized as incremental **Delta** tables on **Azure Databricks** (Spark SQL / PySpark).
- Ran **incremental models** (unique_key + lookback MERGE) in production so daily candidate/exam changes load without full refresh; retained history for restatements.
- Designed **Candidate 360** (grain: 1 row per candidate) with registration, exam stages, status transitions, attempts, and results — fact + dimension model.
- Gated releases with dbt tests (unique, not_null, relationships) in **GitLab CI/CD** so bad keys/stages never reach gold or dashboards.
- Shipped production **Apache Superset** views: candidate 360, stage-funnel drop-off, and freshness monitoring for ops.

Enterprise ICDB Data Platform — Phase 1 **PySpark** · **Databricks** · **SQL** · **CI/CD**

- Built analytics-ready ICDB tables on **Azure Databricks** with **PySpark** / **Spark SQL** and Delta; partitioned for prune-friendly reads by downstream jobs.
- Tuned production SQL/Spark jobs (joins, filters, file layout) using query plans to cut heavy scans and shuffle-heavy patterns.
- Automated environment promotions of data objects via **GitLab CI/CD** with rollback-safe releases.
- Documented **lineage** from warehouse tables to consumers; mapped sources → ICDB and validated grain and business rules in each release.

Enterprise ICDB Data Platform — Phase 2 (Production pipelines) **Integrations** · **SLA** · **Monitoring**

- Operated production data-integration pipelines moving batches between source systems and downstream platforms under SLA.
- Owned monitoring, log analysis, incident RCA, and runbooks so pipeline failures were detected, replayed, and documented.

Production ETL & Warehouse Support **ODI** · **SQL** · **Data quality**

- Owned production ETL (ODI) and warehouse loads: incidents and change requests within SLA; traced bad numbers to joins, filters, and source gaps.
- Fixed mappings, load windows, and row-count checks; mentored juniors on reconciliation, SQL backtracing, and CI/CD for data changes.

Multi-Region Master Data Migration (Wave rollout) **ETL** · **Python** · **Reconciliation**

- Executed wave-based **master-data migrations** across regions with key integrity and completeness gates before cutover.
- Built ETL packages plus SQL/Python reconciliation (row counts, keys, attributes); closed mapping and data-quality gaps before go-live.

Production CDC & Rollback (Database upgrade) **GoldenGate** · **CDC** · **Python**

- Designed production **CDC** and rollback using **Oracle GoldenGate** (Extract/Replicat, INIT/sync, hub/client) plus Data Guard for HA/fallback during a major database upgrade.
- Owned replication lag, trail-file troubleshooting, and abend recovery so capture stayed within cutover limits.
- Wrote Python diff jobs for large extracts; delivered Oracle → MariaDB transforms with automated validation in GitLab CI/CD.

EDUCATION

B.Tech, Electronics and Telecommunication Engineering — Chandigarh Engineering College, Chandigarh 2021 · CGPA 7.8/10